



L2: identity Hessian vs. a gradient measured with all experts scaled to 0.95

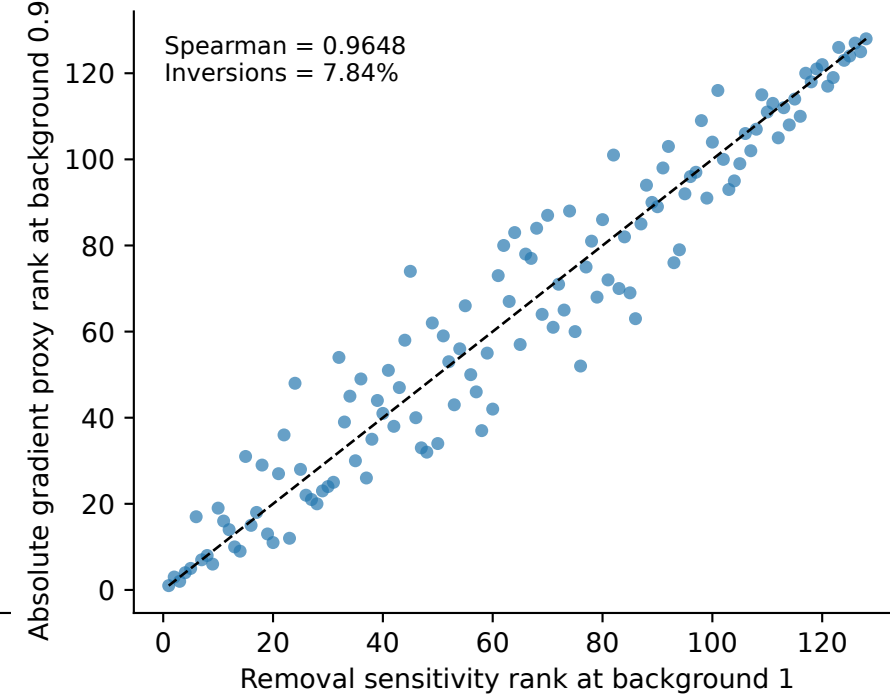
(a) L2: all 128 experts



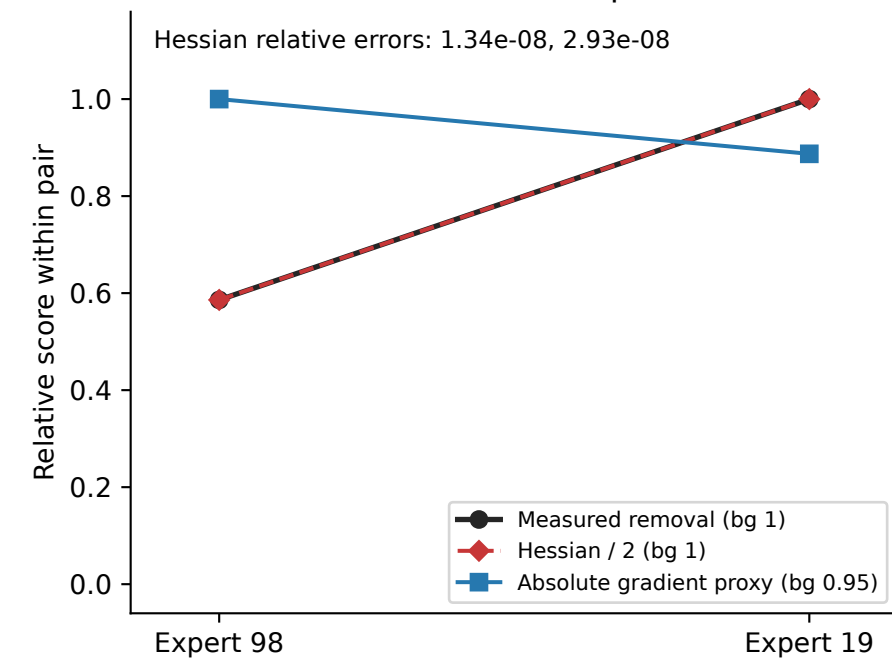
(b) Nonzero first order (signed)



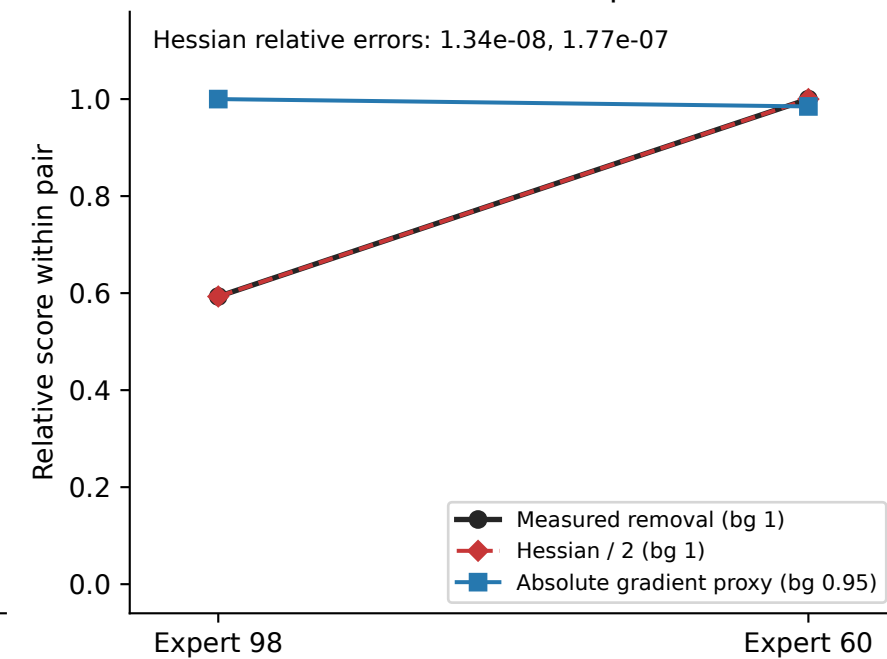
(c) Nonzero first order (absolute)



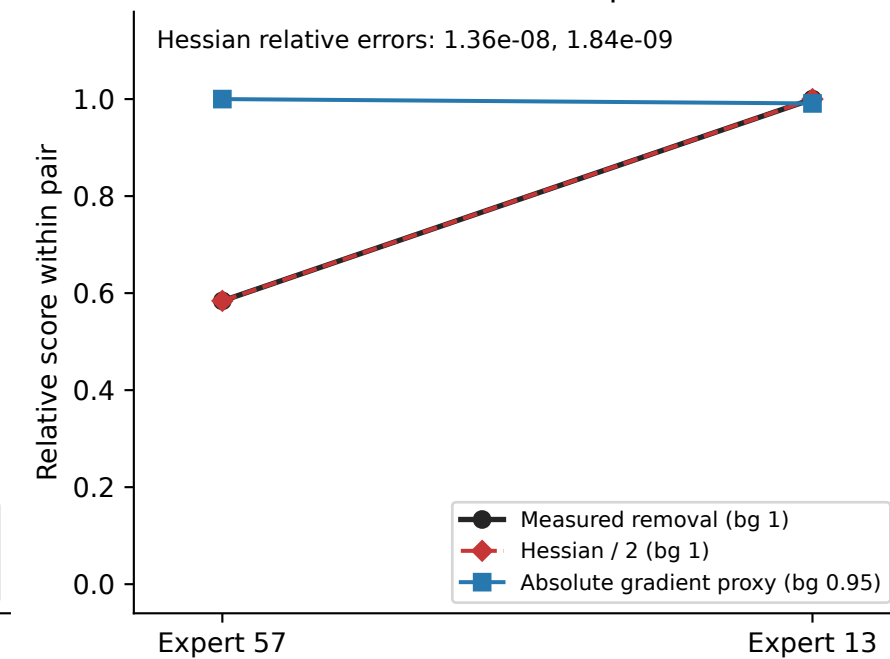
(d) Counterexample 1



(e) Counterexample 2



(f) Counterexample 3



Pair panels: loss and Hessian share the maximum measured loss; the gradient proxy uses its own maximum magnitude.